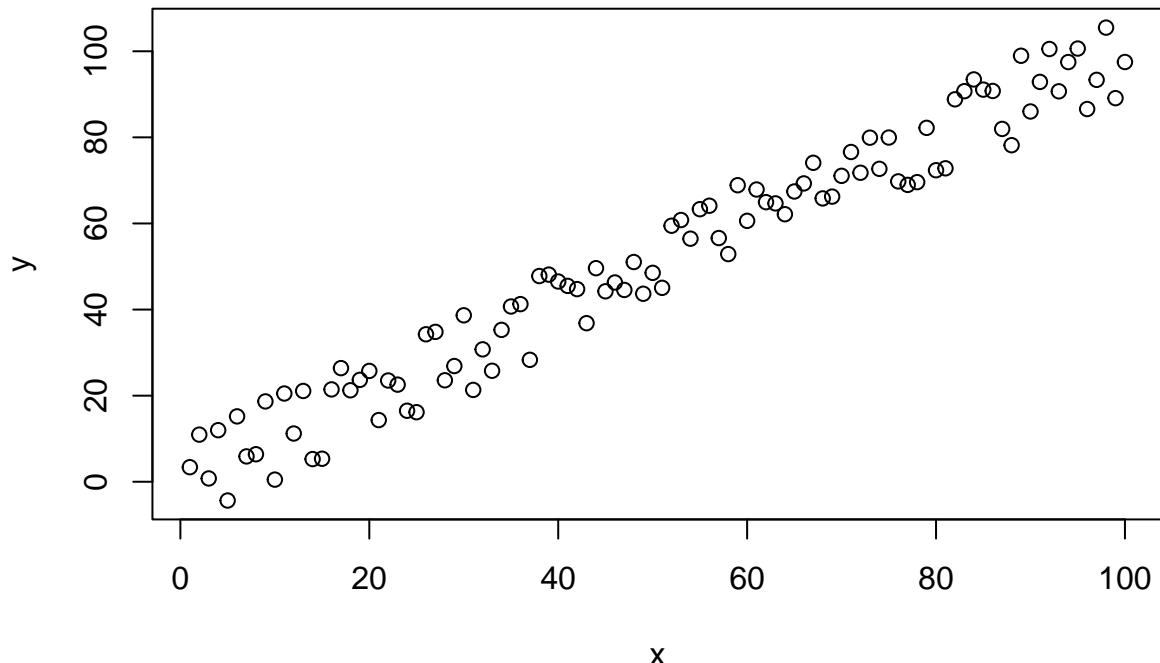


Interrupted Time Series

Let's create some data.

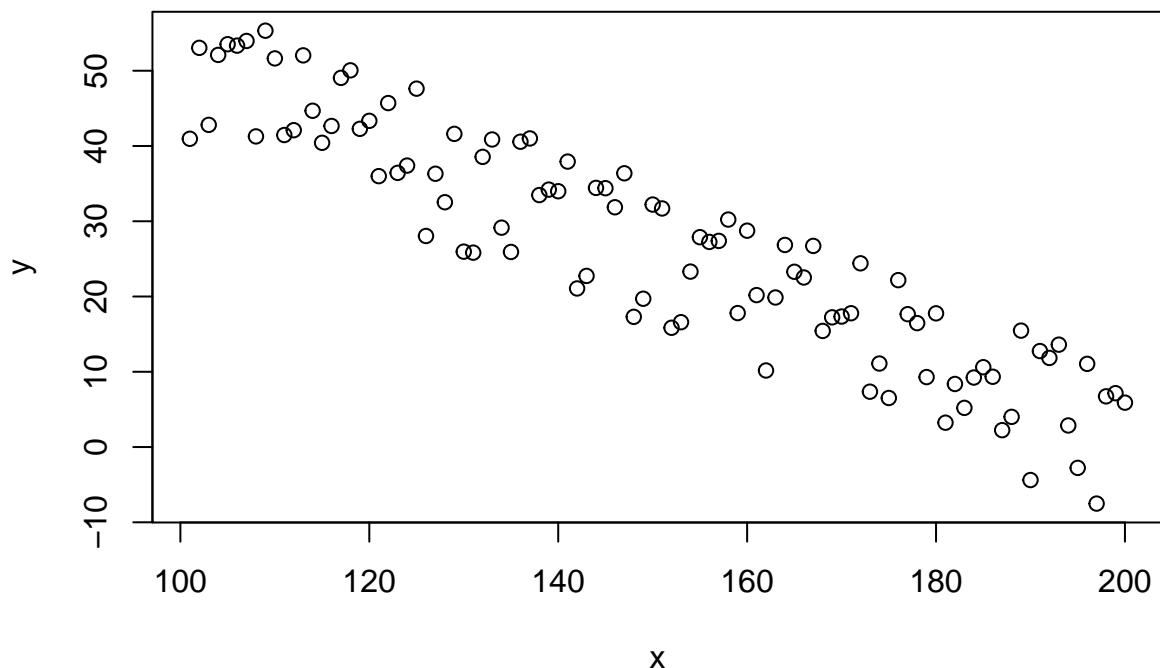
Here's a positive relationship.

```
j = 50  
  
a = data.frame(x=1:100, y=jitter(1:100, j))  
plot(a)
```



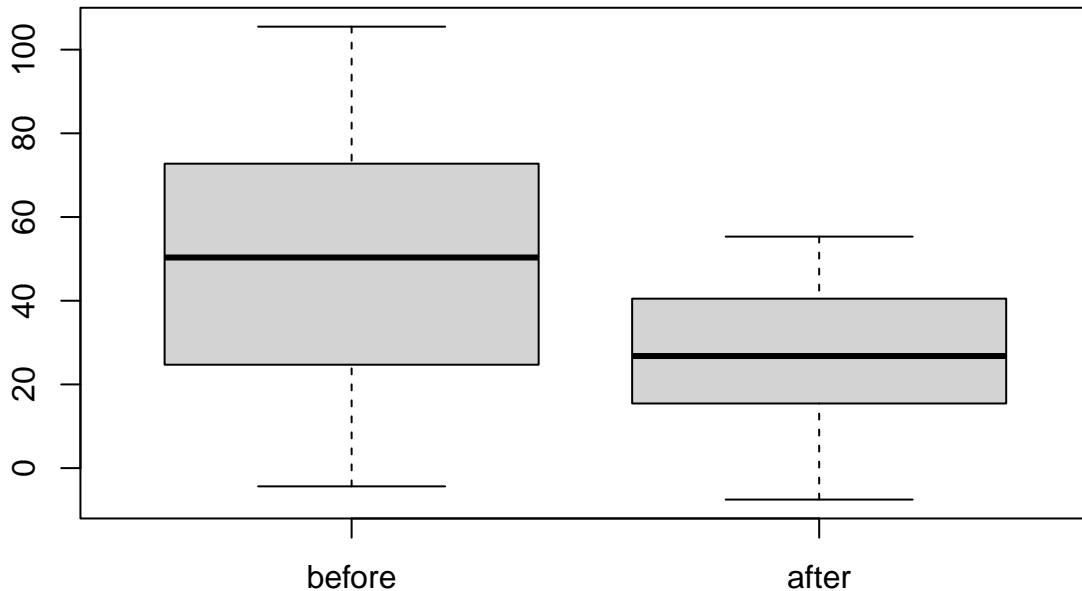
Here's a negative relationship.

```
b = data.frame(x=101:200, y=jitter(100:1, j))  
bb = data.frame(x=101:200, y=jitter(seq(50,0.5,-0.5), 100))  
plot(bb)
```



Are these any different?

```
boxplot(list(before=a$y,after=bb$y))
```



```
t.test(a$y,b$y)
```

```
##  
## Welch Two Sample t-test  
##  
## data: a$y and b$y  
## t = 0.18863, df = 197.99, p-value = 0.8506  
## alternative hypothesis: true difference in means is not equal to 0  
## 95 percent confidence interval:  
## -7.417748 8.986885  
## sample estimates:
```

```

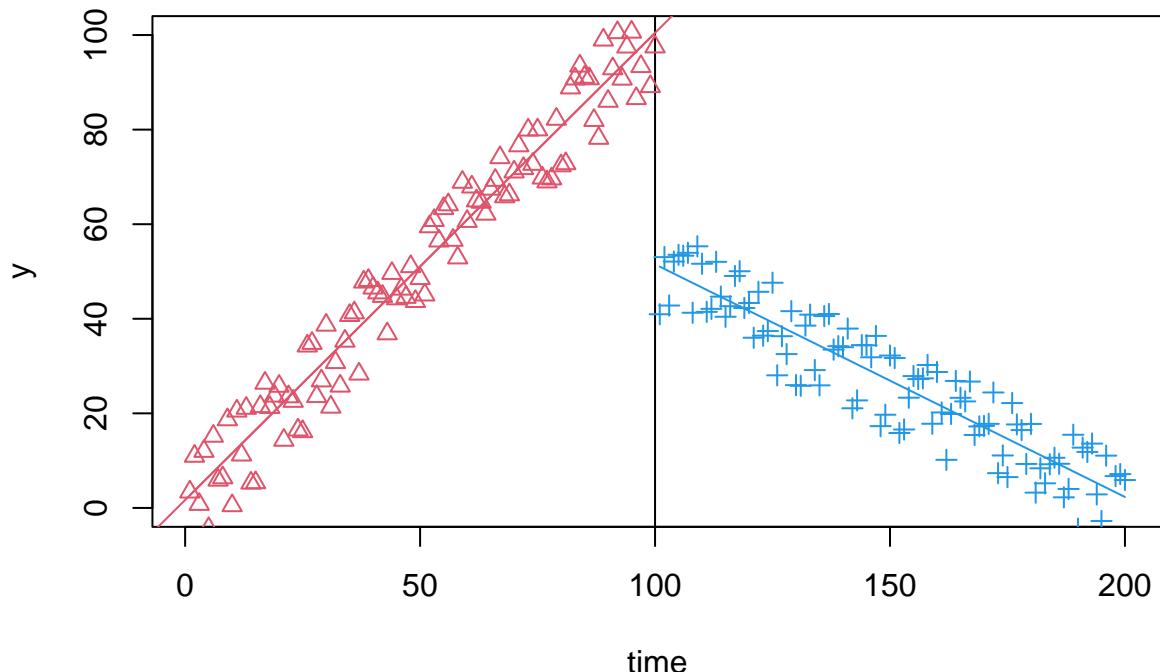
## mean of x mean of y
## 51.56097 50.77640

Let's display them side by side.

plot(x=1:200, y=rep(1,200), type="n", ylim=c(0,100),
      xlab="time", ylab="y")
abline(v=100)
points(a$x, a$y, pch=2, col=2)
points(bb$x, bb$y, pch=3, col=4)

abline(lm(y~x, data=a), col=2)
lines(x=1:100, y=lm(y~x, data=a)$fit, col=2)
# abline(lm(y~x, data=bb), col=4)
lines(x=101:200, y=lm(y~x, data=bb)$fit, col=4)

```



Let's simulate a change in level.

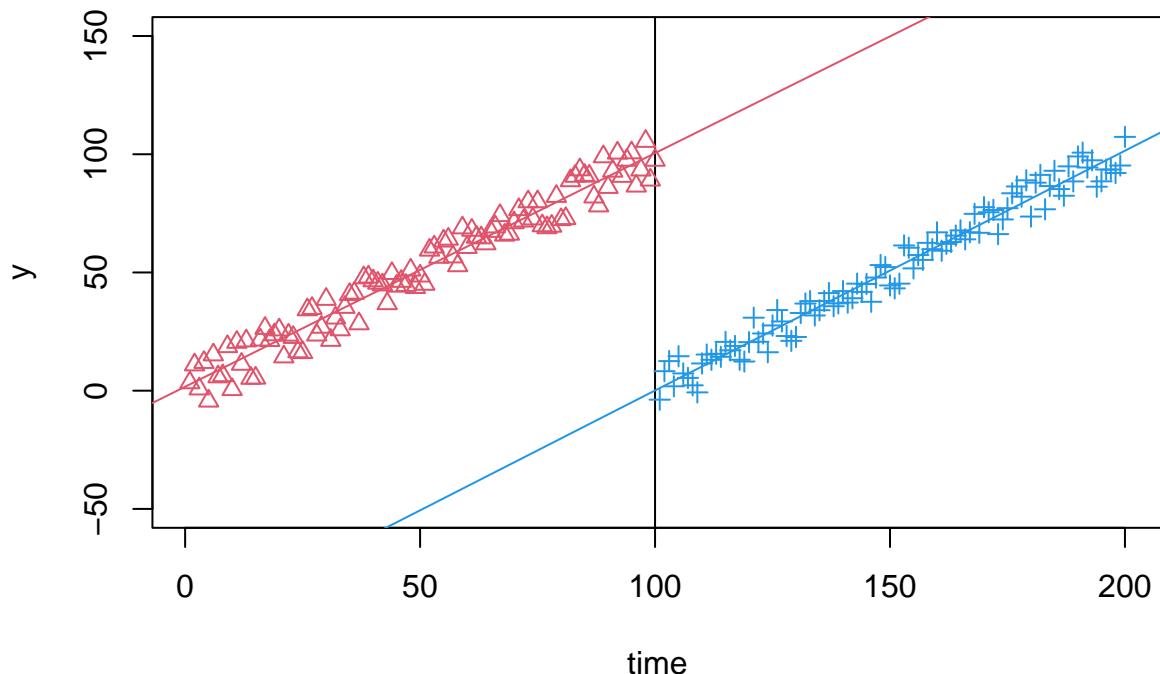
```

a2 = data.frame(x=101:200, y=jitter(1:100, j))

plot(x=1:200, y=rep(1,200), type="n", ylim=c(-50,150),
      xlab="time", ylab="y")
abline(v=100)
points(a$x, a$y, pch=2, col=2)
points(a2$x, a2$y, pch=3, col=4)

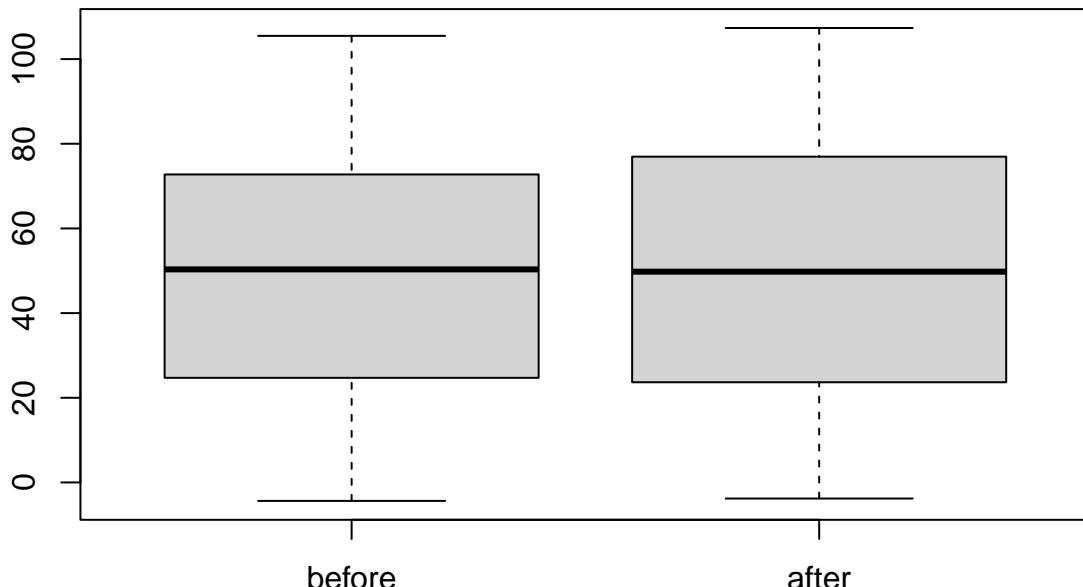
abline(lm(y~x, data=a), col=2)
abline(lm(y~x, data=a2), col=4)

```



We can't capture that with a simple test.

```
boxplot(list(before=a$y, after=a2$y))
```



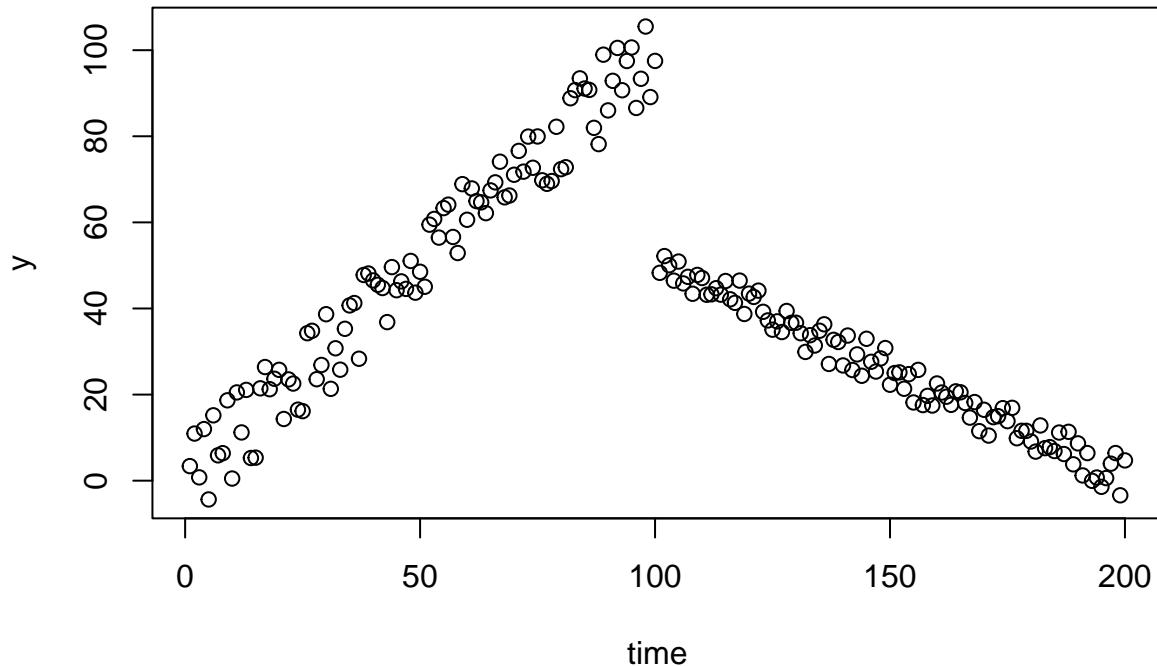
```
t.test(a$y, a2$y)
```

```
##
##  Welch Two Sample t-test
##
## data: a$y and a2$y
## t = 0.062368, df = 197.93, p-value = 0.9503
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -7.992779  8.514855
```

```
## sample estimates:
## mean of x mean of y
## 51.56097 51.29993
```

Now let's go back to the previous example:

```
m = rbind(a, data.frame(x=101:200, y=jitter(seq(50,0.5,-0.5), j)))
plot(m$x, m$y, xlab="time", ylab="y")
```



Here's what a simple model might look like:

```
summary(lm(y~x, data=m))

##
## Call:
## lm(formula = y ~ x, data = m)
##
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -55.846 -17.459  -1.771  11.772  66.756 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 52.17421   3.62637 14.387 < 2e-16 ***
## x          -0.13715   0.03129 -4.383 1.89e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 25.55 on 198 degrees of freedom
## Multiple R-squared:  0.08846,    Adjusted R-squared:  0.08385 
## F-statistic: 19.21 on 1 and 198 DF,  p-value: 1.895e-05
```

Let's see if we can model those trends and change in level explicitly.


```

## 
## Call:
## lm(formula = y ~ time + intervention + time_after_intervention,
##      data = m)
## 
## Residuals:
##       Min     1Q   Median     3Q    Max 
## -11.1442 -2.9659  0.2115  3.4002  9.3657 
## 
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|)    
## (Intercept)             1.67710   0.97135   1.727   0.0858 .  
## time                  0.98780   0.01670  59.153 <2e-16 *** 
## interventionTRUE      -49.75945   1.36350 -36.494 <2e-16 *** 
## time_after_intervention -1.49229   0.02362 -63.190 <2e-16 *** 
## --- 
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1 
## 
## Residual standard error: 4.82 on 196 degrees of freedom
## Multiple R-squared:  0.9679, Adjusted R-squared:  0.9674 
## F-statistic:  1968 on 3 and 196 DF,  p-value: < 2.2e-16

```

Q: Can you achieve the same result (i.e., capture both trends and the change in level) with only two variables?
A: Yes, with an interaction term!

```
rdd2 = lm(y ~ time * intervention, data=m)
summary(rdd2)
```

```

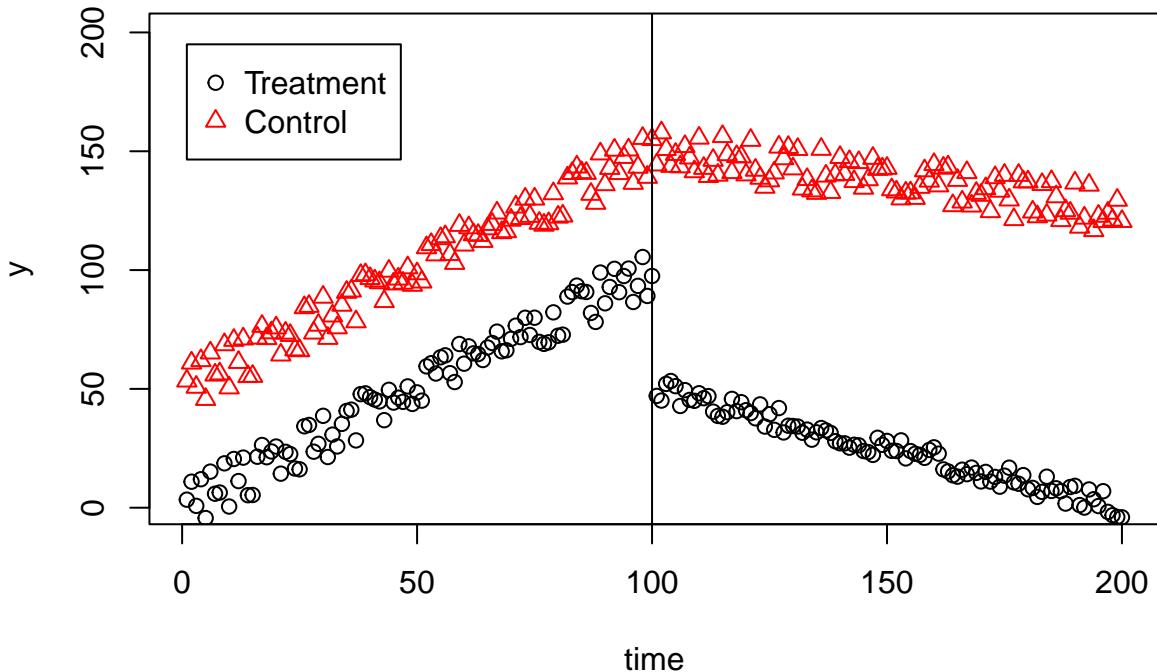
## 
## Call:
## lm(formula = y ~ time * intervention, data = m)
## 
## Residuals:
##       Min     1Q   Median     3Q    Max 
## -11.1442 -2.9659  0.2115  3.4002  9.3657 
## 
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|)    
## (Intercept)             1.67710   0.97135   1.727   0.0858 .  
## time                  0.98780   0.01670  59.153 <2e-16 *** 
## interventionTRUE      99.46970   2.73716  36.340 <2e-16 *** 
## time:interventionTRUE -1.49229   0.02362 -63.190 <2e-16 *** 
## --- 
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1 
## 
## Residual standard error: 4.82 on 196 degrees of freedom
## Multiple R-squared:  0.9679, Adjusted R-squared:  0.9674 
## F-statistic:  1968 on 3 and 196 DF,  p-value: < 2.2e-16

```

Now let's add a control series.

```
a2 = a
names(a2) = c("x", "yt")
df = rbind(a2, data.frame(x=101:200, yt=jitter(seq(50,0.5,-0.5), j)))
df$yc = jitter(50) + df$yt
df[df$x>=100,]$yc = jitter(seq(150,125,-0.25), 4*j)
```

```
{
  plot(df$x, type="n", xlab="time", ylab="y")
  points(df$x, df$yt)
  points(df$x, df$yc, col = "red", pch=2)
  legend(1, 195, legend=c("Treatment", "Control"),
         col=c("black", "red"), pch=c(21,2))
  abline(v=100)
}
```



And set up the ITS variables.

```
dfm = data.frame(time = df$x, y = c(df[c("x","yt")]$yt,df[c("x","yc")]$yc))

dfm$group = c(rep("treated",200), rep("control",200))
dfm$intervention = dfm$time > 100
dfm$time_after_intervention = ifelse(dfm$time > 100, dfm$time - 100, 0)

rdd2c = lm(y ~ time
            + intervention
            + time_after_intervention
            + group
            + group:time
            + group:intervention
            + group:time_after_intervention
            , data=dfm)
summary(rdd2c)

## 
## Call:
## lm(formula = y ~ time + intervention + time_after_intervention +
##     group + group:time + group:intervention + group:time_after_intervention,
##     data = dfm)
## 
```

```

## Residuals:
##      Min       1Q   Median      3Q      Max
## -11.1442  -3.7472 -0.1694  4.4578  9.9158
##
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|)
## (Intercept)                51.481653   1.122071  45.881 < 2e-16 ***
## time                      0.992247   0.019290  51.438 < 2e-16 ***
## interventionTRUE          -0.638389   1.575079 -0.405   0.685
## time_after_intervention    -1.243818   0.027280 -45.594 < 2e-16 ***
## grouptreated               -49.804550   1.586847 -31.386 < 2e-16 ***
## time:grouptreated          -0.004448   0.027280 -0.163   0.871
## interventionTRUE:grouptreated -49.764833   2.227498 -22.341 < 2e-16 ***
## time_after_intervention:grouptreated -0.253254   0.038580 -6.564 1.66e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5.568 on 392 degrees of freedom
## Multiple R-squared:  0.9874, Adjusted R-squared:  0.9872
## F-statistic:  4404 on 7 and 392 DF,  p-value: < 2.2e-16

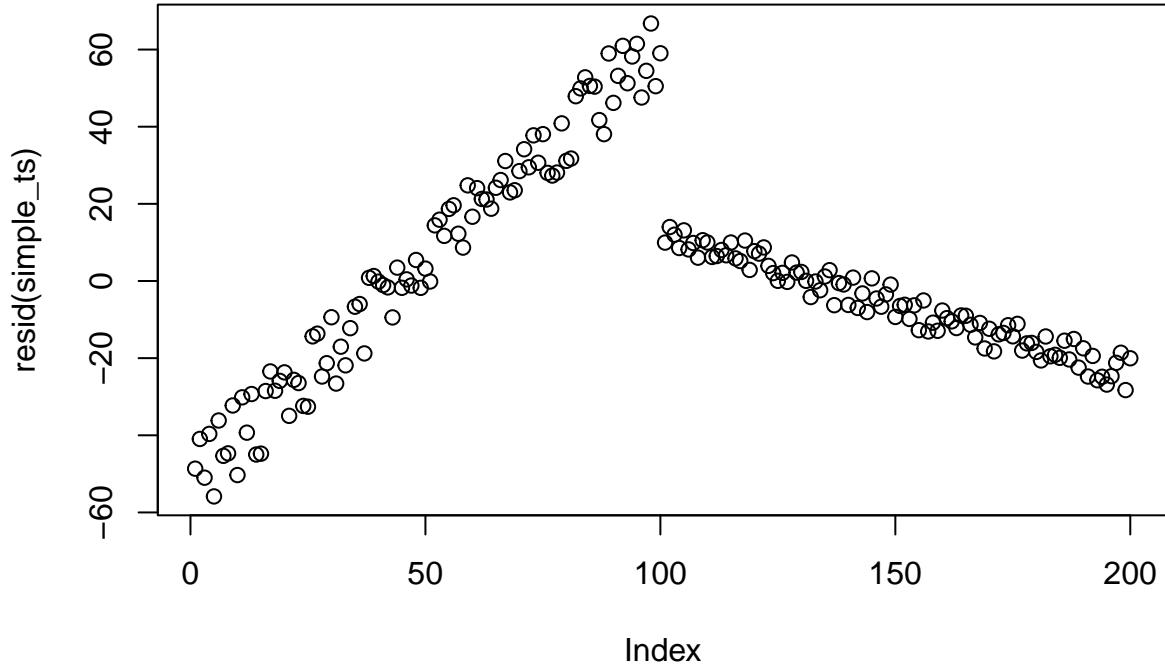
```

Is there autocorrelation?

```

simple_ts = lm(y ~ time, data=m)
plot(resid(simple_ts))

```

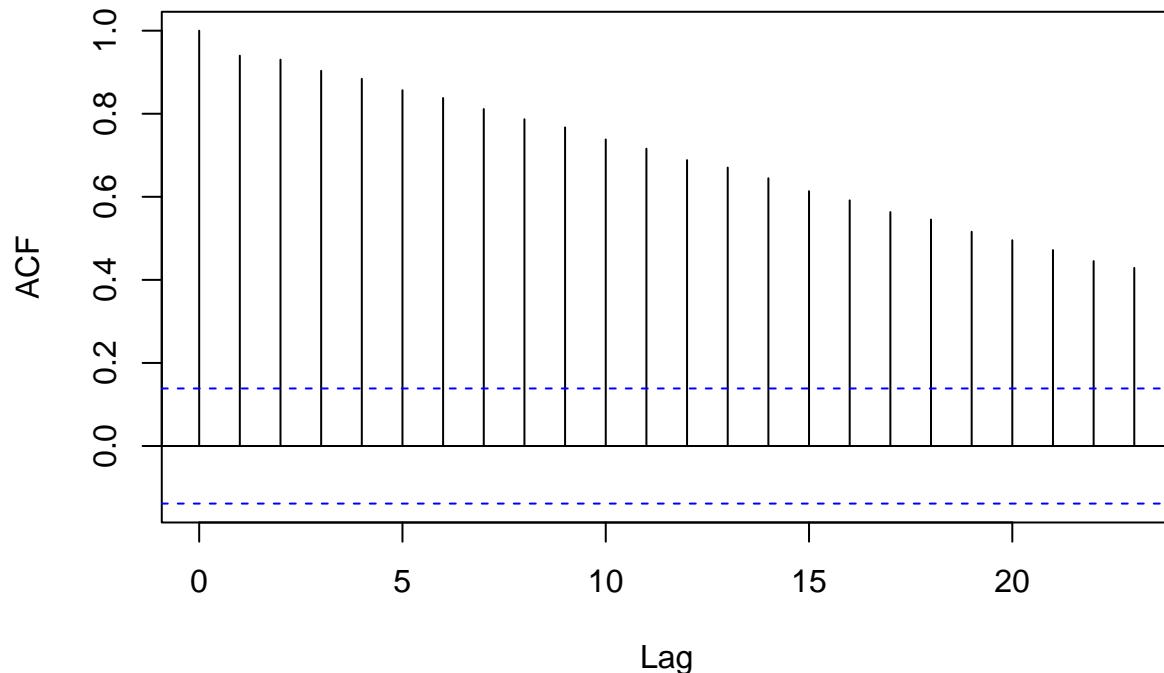


```

# alternatively
acf(resid(simple_ts))

```

Series resid(simple_ts)



To formally test for autocorrelation, we can use the Durbin-Watson test

```
library(lmtest)

## Warning: package 'lmtest' was built under R version 4.4.1
## Loading required package: zoo
##
## Attaching package: 'zoo'

## The following objects are masked from 'package:base':
##       as.Date, as.Date.numeric
dwtest(m$y ~ m$time)

##
## Durbin-Watson test
##
## data: m$y ~ m$time
## DW = 0.098841, p-value < 2.2e-16
## alternative hypothesis: true autocorrelation is greater than 0
```

From the p-value, we know that there is autocorrelation in the time series

A solution to this problem could be to use more advanced time series analysis (e.g., ARIMA) to adjust for seasonality and other dependency, or to use mixed-effects models when modeling multiple individual “treated” time series jointly.